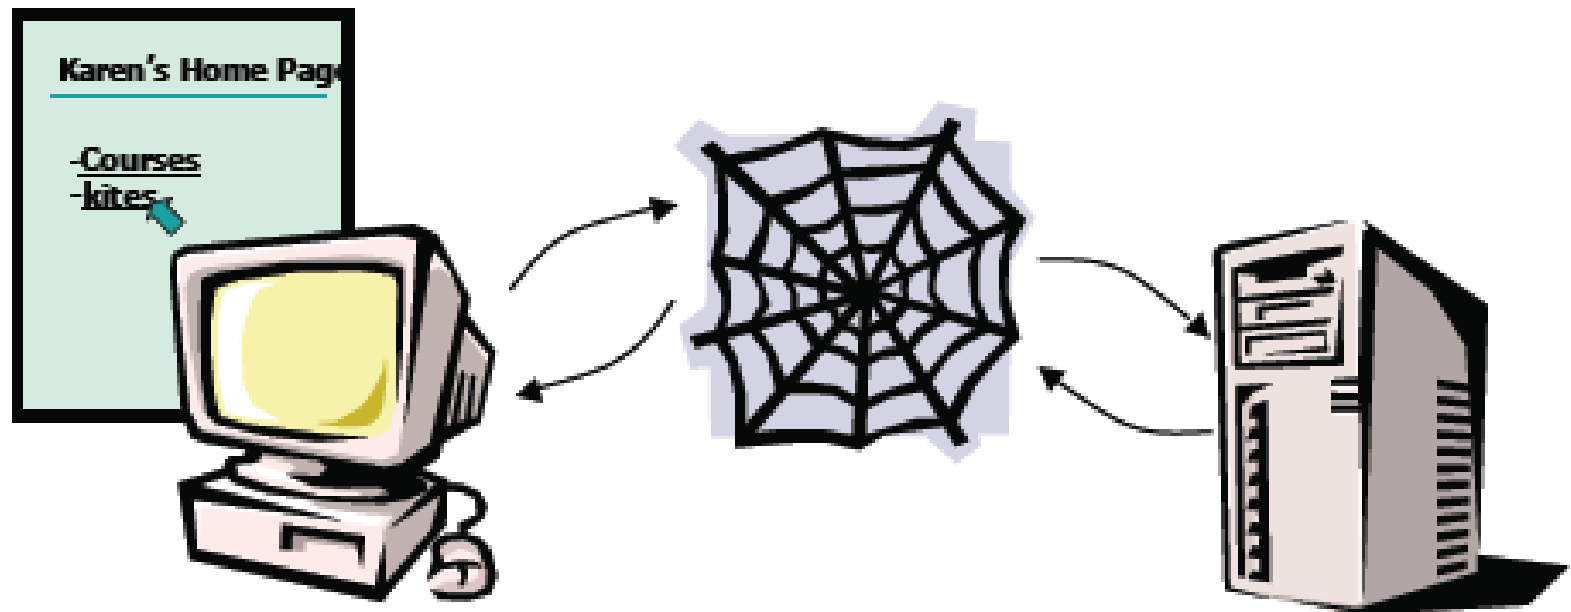




System Programming

Sockets

Simple Web Request



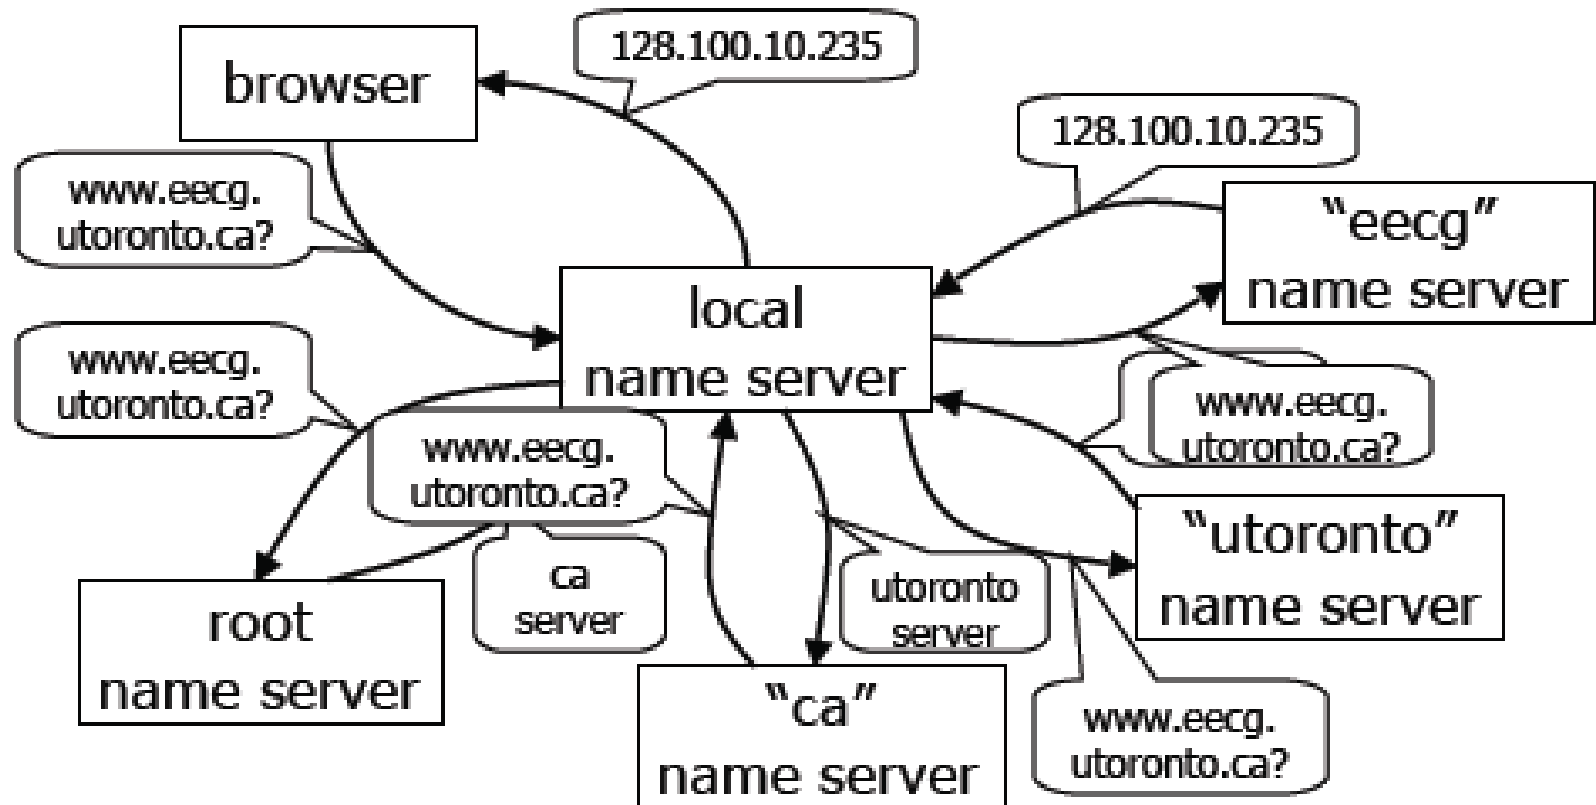


How do we find the server?

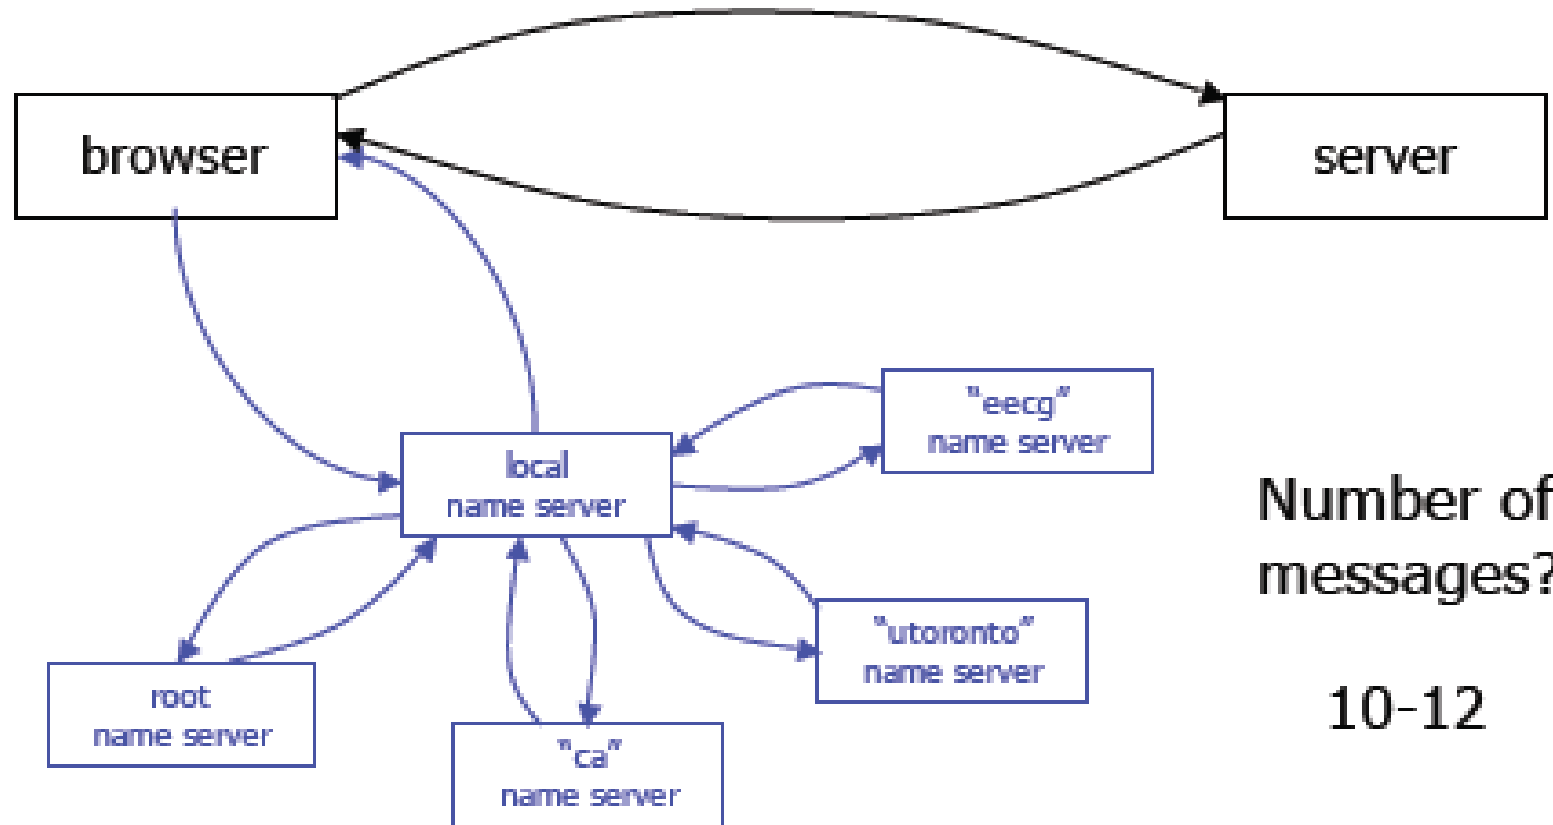
- Every computer on the Internet has an Internet address.
- Called an IP address (Internet Protocol)
- An IP address is 4 8-bit numbers separated by dots.

`www.ktu.edu.tr = 193.140.168.225`

Domain Name Servers



This is getting complicated!



Protocols

Invoice:

Customer: Karen Reid

Order No: 5379

Qty:		Unit Price	Total
1	Athalon	219.00	219.00
2	128 MB	149.95	299.90
	Subtotal		518.90
	Tax		77.84
	TOTAL		596.74

Karen Reid

Feb 18, 2001

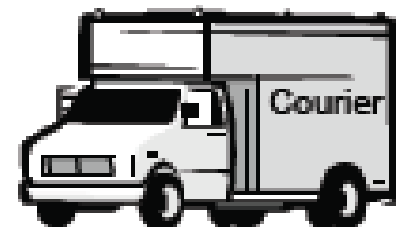
Payable to: CPUS are us \$596.74
Five hundred ninety six 74/100

CPUS are us

Karen Reid
Dept. of Computer Science
University of Toronto

Karen Reid

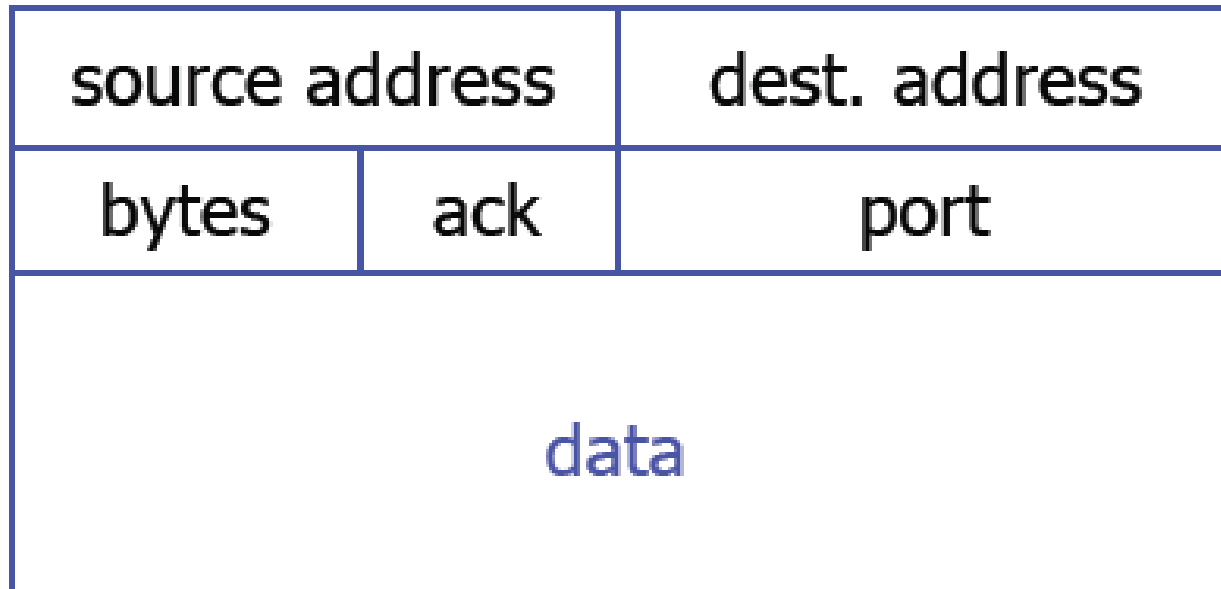
CPUS are us
0 College Street
Toronto Ontario M5S 3G4



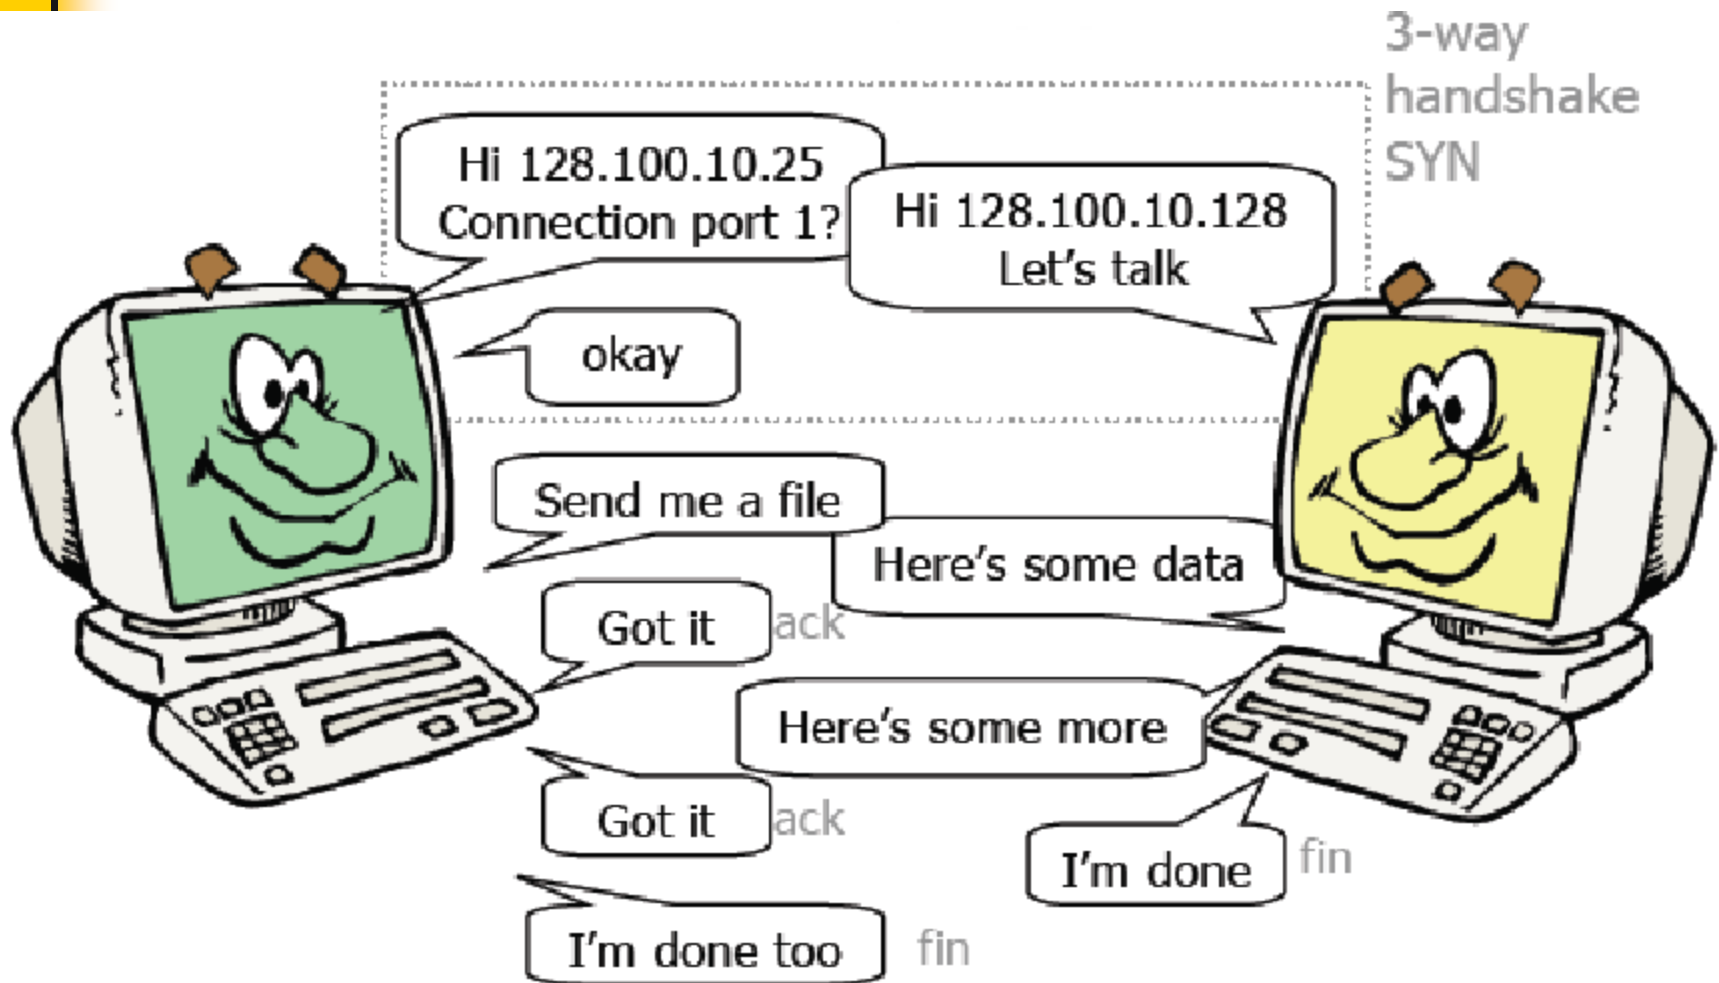


TCP/IP

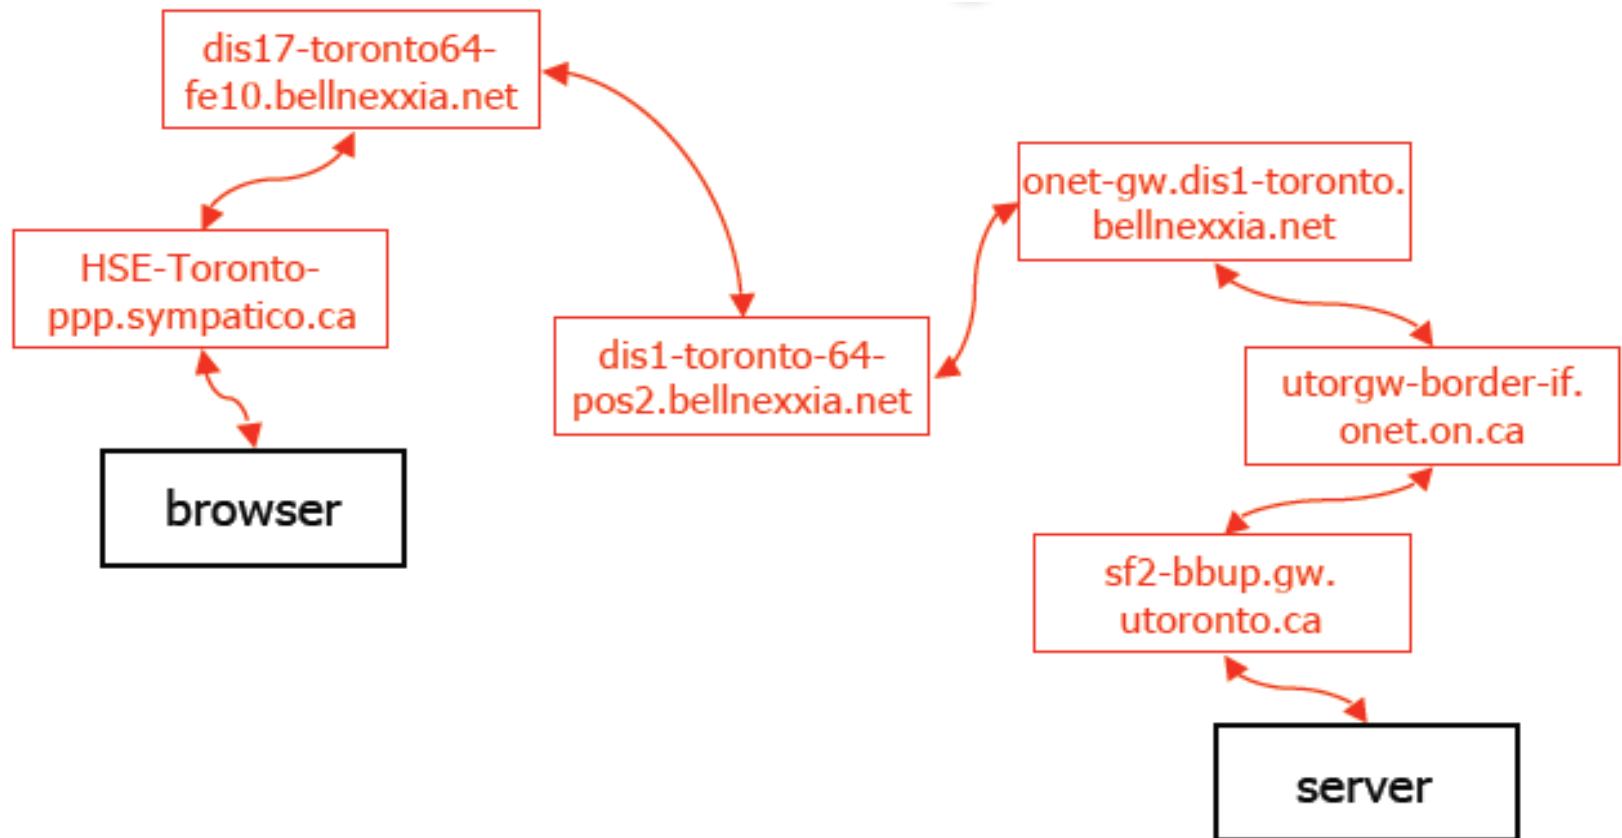
- Transmission Control Protocol.
- Tells us how to package up the data.



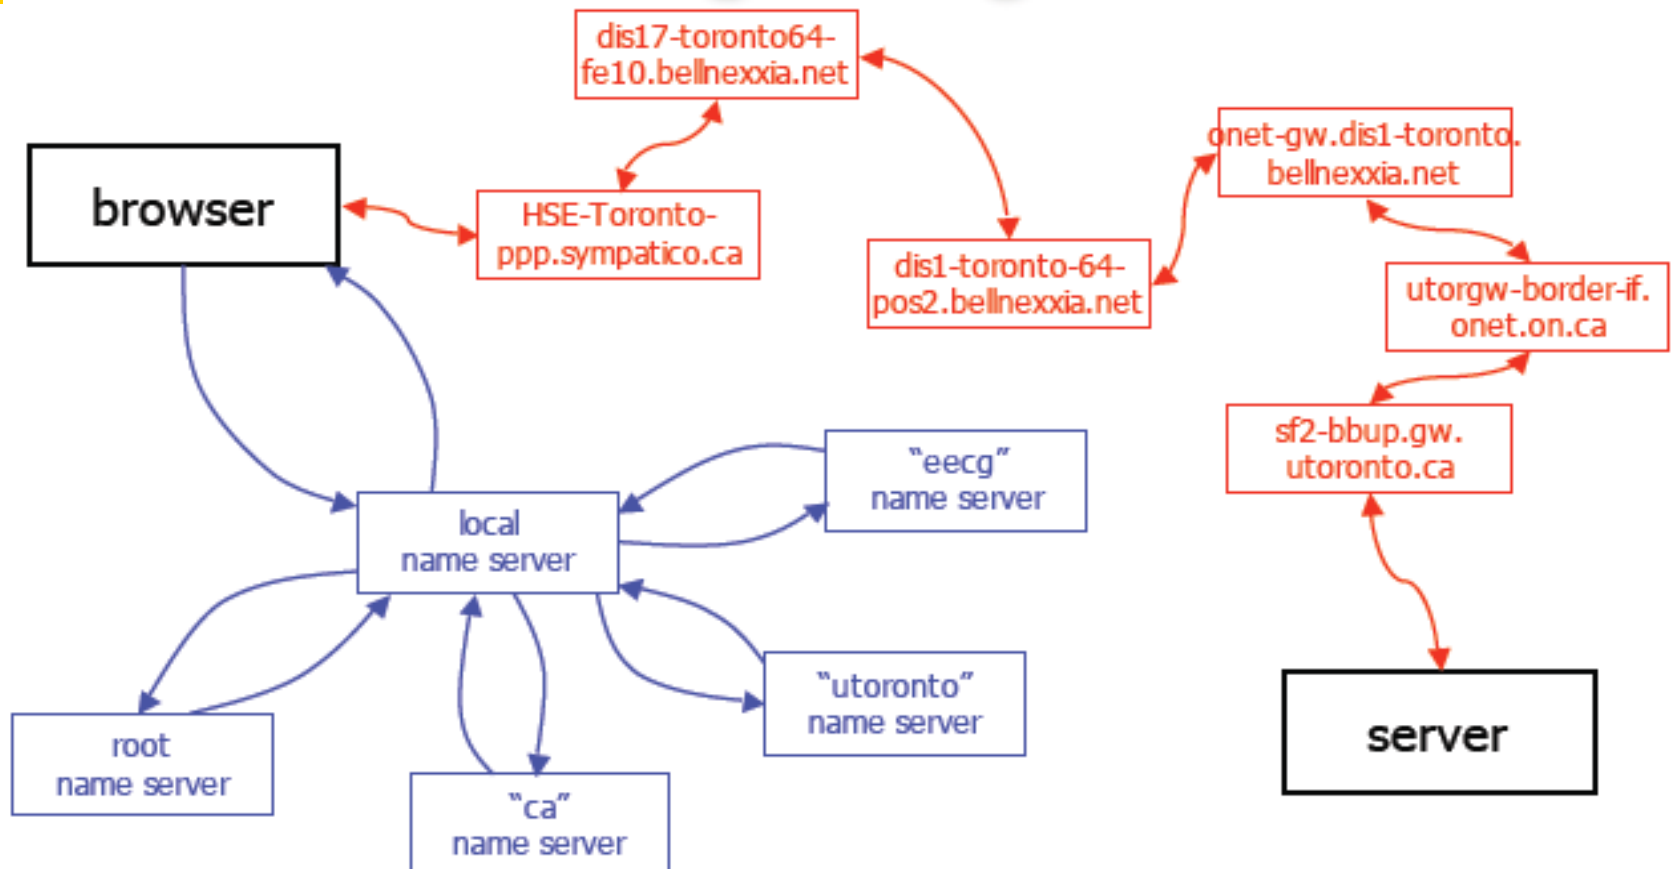
TCP Connection



Routing



Putting it together





How many messages?

- It depends on the size of the web page we retrieve.
- If the web page is 75 Kbytes (small!) it will be broken up into 103 IP packets.
- Remember DNS took 10 messages

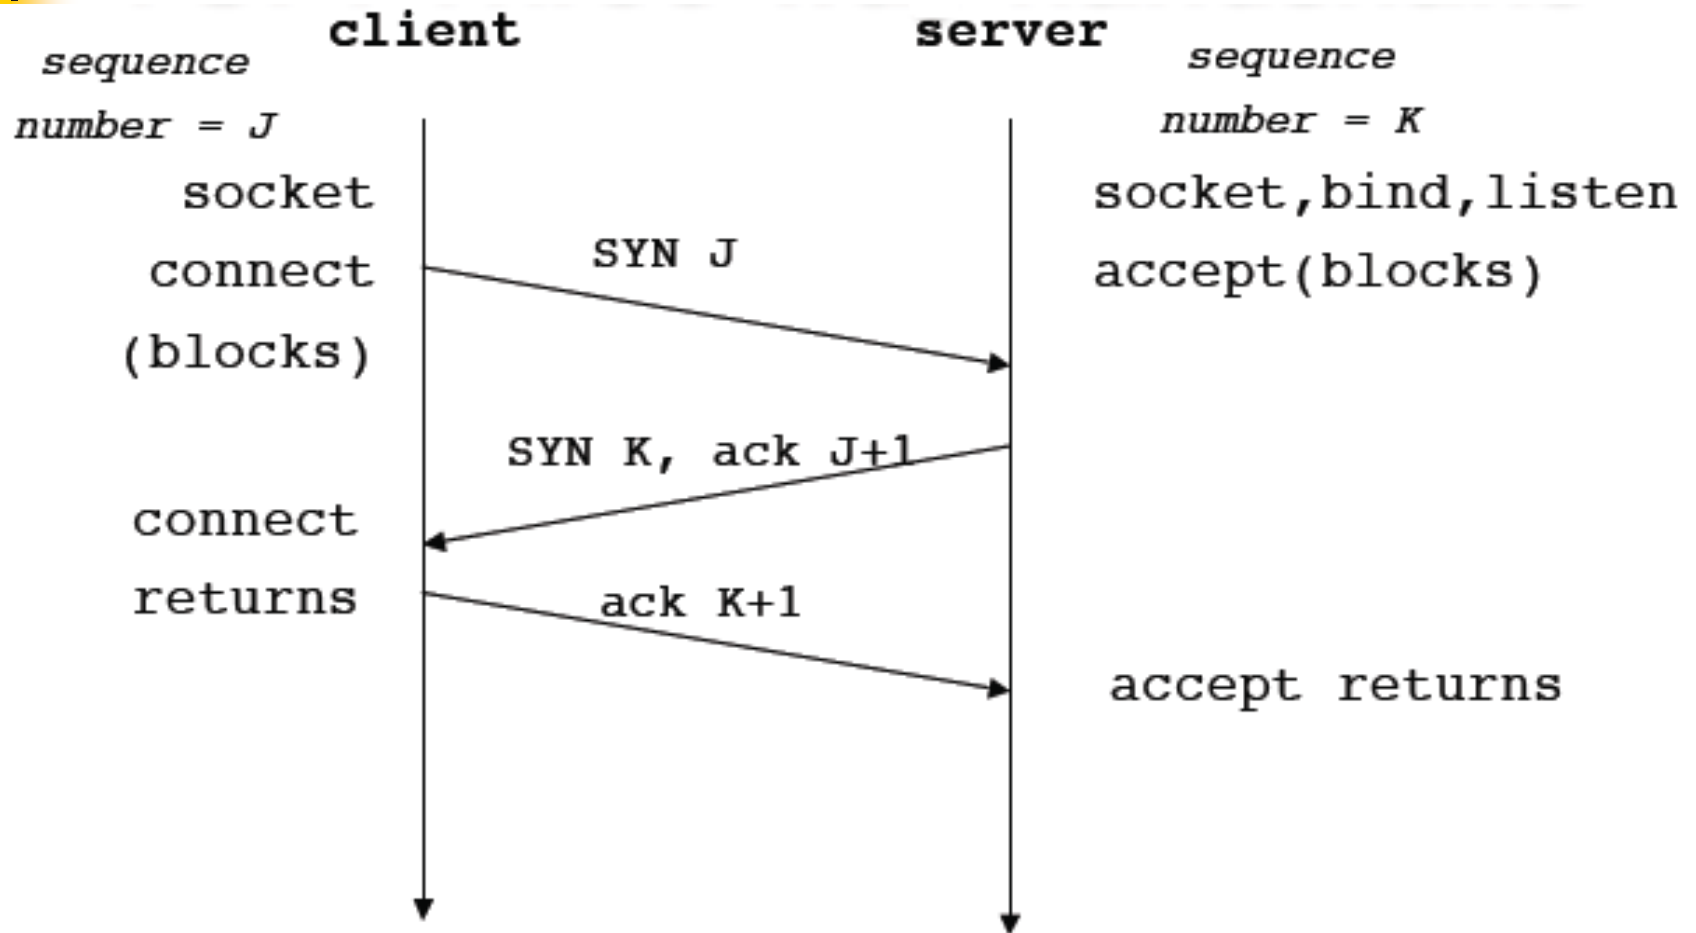
$$10 + 103 \times 7 \text{ hops} = 731 \text{ messages!}$$



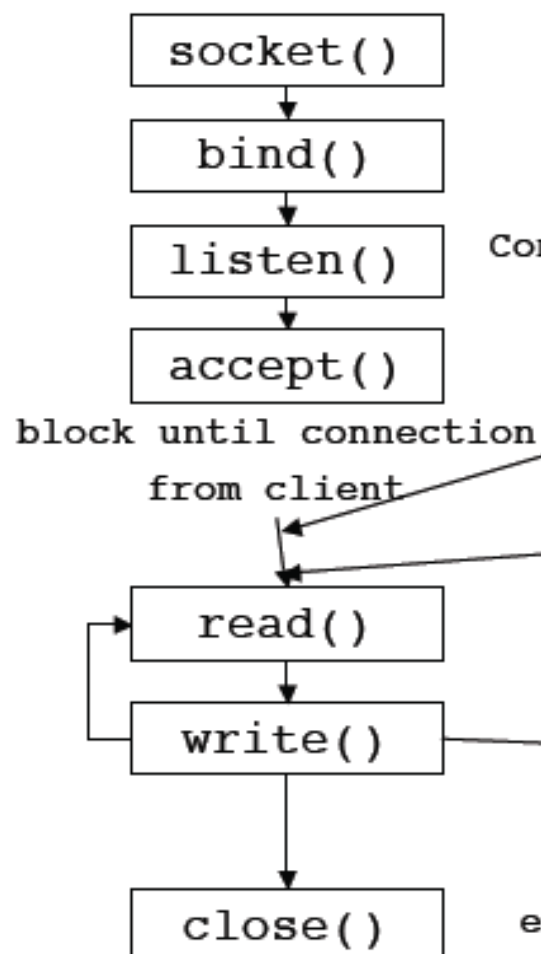
Sockets

- One form of communication between processes.
- Similar to pipes, except sockets can be used between processes on different machines.
- Use file descriptors to refer to sockets.
- Built on top of TCP layer

TCP: Three-way handshake



TCP Server

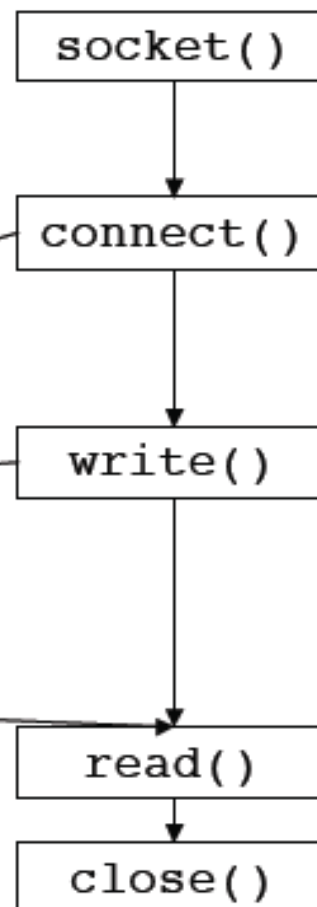


TCP Client

Connection establishment
(3-way handshake)

data transfer

end-of-file notification





Connection-Oriented

■ Server

- Create a socket: `socket ()`
- Assign a name to a socket: `bind ()`
- Establish a queue for connections:
`listen ()`
- Get a connection from the queue: `accept ()`

■ Client

- Create a socket: `socket ()`
- Initiate a connection: `connect ()`



Socket Types

- Two main categories of sockets
 - UNIX domain: both processes on the same machine
 - INET domain: processes on different machines
- Three main types of sockets:
 - SOCK_STREAM: the one we will use
 - SOCK_DGRAM: for connectionless sockets
 - SOCK_RAW



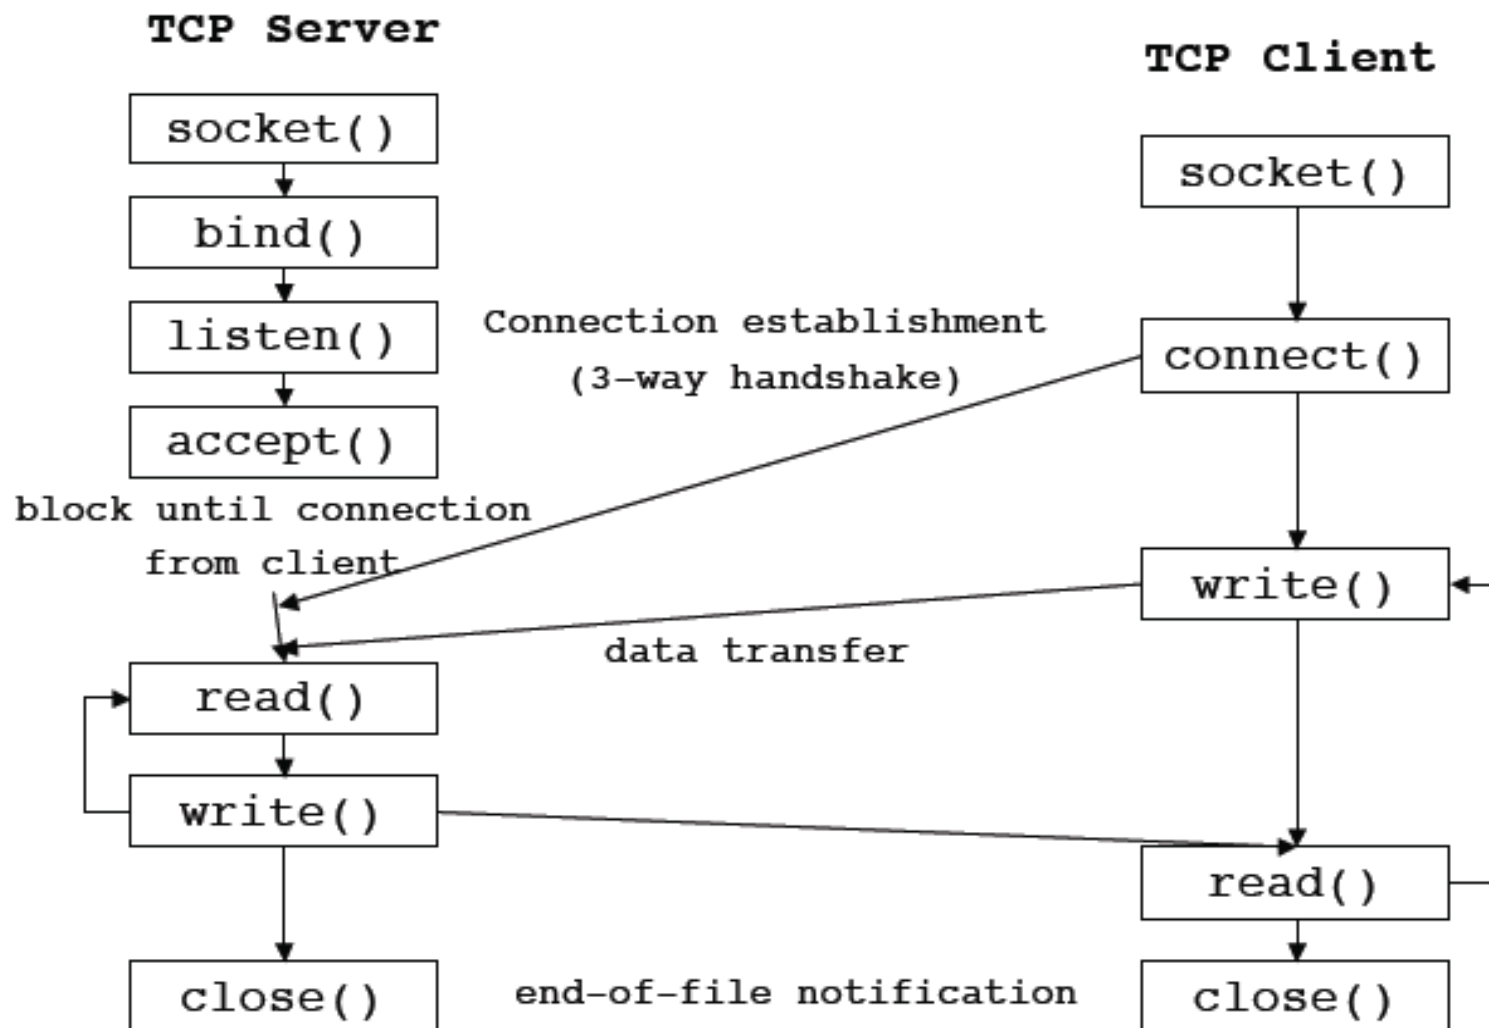
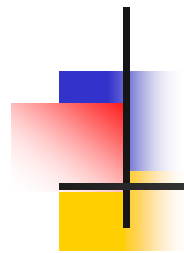
Addresses and Ports

- A **socket pair** is the two endpoints of the connection.
- An endpoint is identified by an **IP address and a port**.
- IPv4 addresses are 4 8-bit numbers:
 - 128.100.31.156 = penguin
 - 128.100.31.4 = eddie
- Ports
 - because multiple processes can communicate with a single machine we need another identifier.



More on Ports

- Well-known ports: 0-1023
 - 80 = web
 - 21 = ftp
 - 22 = ssh
 - 25 = smtp (mail)
 - 23 = telnet
 - 194 = irc
- Registered ports: 1024-49151
 - 2709 = supermon
 - 26000 = quake
- Dynamic (private) ports: 49152-65535
 - You should pick ports in this range to avoid overlap





Server side

```
int socket(int family, int type,  
           int protocol);
```

- family specifies protocol family:
 - PF_INET – IPv4
 - PF_LOCAL – Unix domain
- type
 - SOCK_STREAM, SOCK_DGRAM, SOCK_RAW
- protocol
 - set to 0 except for RAW sockets
- returns a socket descriptor



bind to a name

```
int bind(int sockfd,
        const struct sockaddr *servaddr,
        socklen_t addrlen);
```

- sockfd – returned by socket

```
struct sockaddr_in{
    short    sin_family; /*PF_INET */
    u_short  sin_port;
    struct   in_addr sin_addr;
    char     sin_zero[8];
}
```

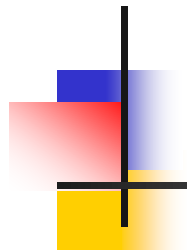
- sin_addr can be set to INADDR_ANY to communicate with any host

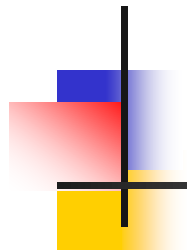


Set up queue in kernel

```
int listen(int sockfd, int backlog)
```

- after calling `listen`, a socket is ready to accept connections
- prepares a queue in the kernel where partially completed connections wait to be accepted.
- `backlog` is the maximum number of partially completed connections that the kernel should queue.







Complete the connection

```
int accept(int sockfd,  
           struct sockaddr *cliaddr,  
           socklen_t *addrlen);
```

- blocks waiting for a connection (from the queue)
- returns a new descriptor which refers to the TCP connection with the client
 - `sockfd` is the listening socket
 - `cliaddr` is the address of the client
- reads and writes on the connection will use the socket returned by `accept`



Client side

- `socket ()` – same as server, to say “how” we are going to talk

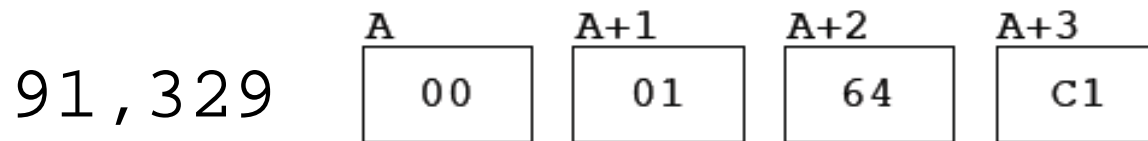
```
int connect(int sockfd,  
            const struct sockaddr *servaddr,  
            socklen_t addrlen);
```

- the kernel will choose a dynamic port and source IP address.
- returns 0 on success and -1 on failure setting `errno`.
- initiates the three-way handshake.

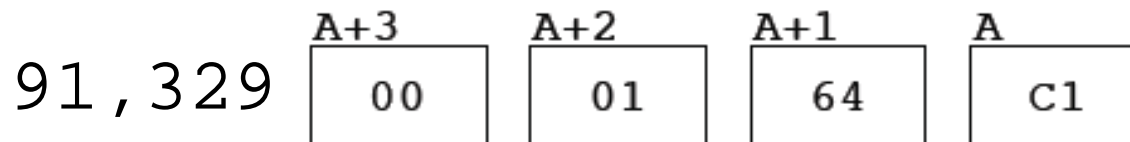


Byte order

- Big-endian



- Little-endian



- Intel is little-endian, and Sparc is big-endian



Network byte order

- To communicate between machines with unknown or different “endian-ness” we convert numbers to network byte order (bigendian) before we send them.
- There are functions provided to do this:
 - `unsigned long htonl(unsigned long)`
 - `unsigned short htons(unsigned short)`
 - `unsigned long ntohl(unsigned long)`
 - `unsigned short ntohs(unsigned short)`